


# CS 1501

## Compression



# ***The Compression "Problem"***

Given a representation of some information, encode that information using fewer bits than the original representation


# Why compress?

- Can get more use out a disk of a given size
- Can get more use out of memory
  - E.g., free up memory by compressing inactive sections
    - Faster than paging
    - Built in to OSX Mavericks and later
- Can reduce the amount data transmitted
  - Faster file transfers
  - Cut power usage on mobile devices

# Approaches to compression

- Can be grouped into two broad categories...

# Lossy Compression



- Information is permanently lost in the compression process
- Examples:
  - MP3, H264, JPEG
- With audio/video files this typically isn't a huge problem as human users might not be able to perceive the difference

# Lossy examples

- MP3
  - “Cuts out” portions of audio that are considered beyond what most people are capable of hearing
- JPEG




40K



28K

# Lossless Compression



- Input can be recovered from compressed data exactly
- Examples:
  - zip files, FLAC

# ***The Lossless Compression "Problem"***

Given a representation of some information, encode the same information exactly using fewer bits than the original representation



# Huffman Compression

- Works on arbitrary bit strings, but pretty easily explained using characters
- Consider the ASCII character set
  - Essentially blocks of codes
    - In general, to fit  $R$  potential characters in a block, you need  $\lg R$  bits of storage per block
      - Consequently,  $n$  bits storage blocks represent  $2^n$  characters
    - Each 8 bit code block represents one of 256 possible characters in extended ASCII
    - Easy to encode/decode

# Considerations for compressing ASCII

- What if we used *variable length* codewords instead of the constant 8? Could we store the same info in less space?
  - Different characters are represented using codes of different bit lengths
  - If all characters in the alphabet have the same usage frequency, we can't beat block storage
    - On a character by character basis...
  - What about different usage frequencies between characters?
    - In English, R, S, T, L, N, E are used much more than Q or X

# Variable length encoding


- Decoding was easy for block codes
  - Grab the next 8 bits in the bitstring
  - How can we decode a bitstring that is made of of variable length code words?
  - BAD example of variable length encoding:

1	A
00	T
01	K
001	U
100	R
101	C
10101	N

# Variable length encoding for lossless compression

- Codes must be *prefix free*
  - No code can be a prefix of any other in the scheme
  - Using this, we can achieve compression by:
    - Using fewer bits to represent more common characters
    - Using longer codes to represent less common characters

# How can we create these prefix-free codes?

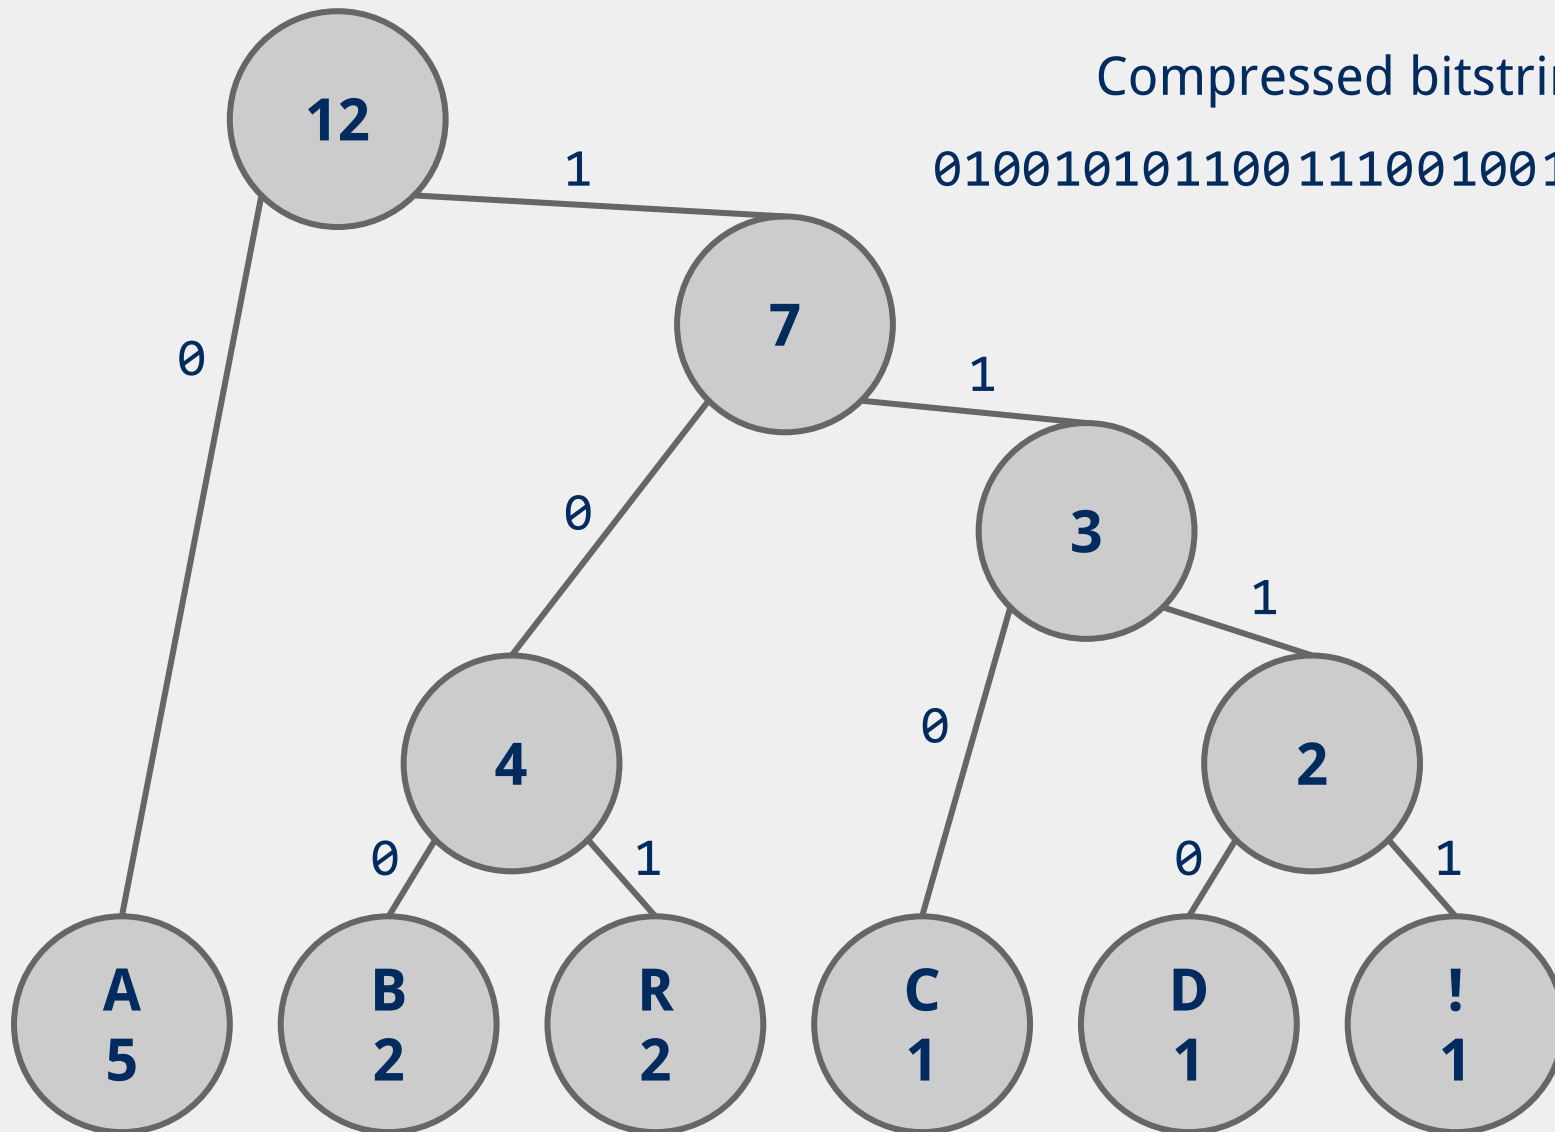


Huffman encoding!

# Generating Huffman codes

- Assume we have  $K$  characters that are used in the file to be compressed and each has a weight (its frequency of use)
- Create a forest,  $F$ , of  $K$  single-node trees, one for each character, with the single node storing that char's weight
- while more than 1 tree in the  $F$ :
  - Select  $T1, T2 \in F$  that have the smallest weights in  $F$
  - Create a new tree node  $N$  whose weight is the sum of  $T1$  and  $T2$ 's weights
  - Remove  $T1$  and  $T2$  from  $F$
  - Add  $T1$  and  $T2$  as children (subtrees) of  $N$
  - Add the new tree rooted by  $N$  to  $F$
- Build a tree for "ABRACADABRA!"

# ABRACADABRA!



Compressed bitstring:

0100101011001110010010101111

# Implementation concerns

- Need to efficiently be able to select lowest weight trees to merge when constructing the trie
  - Can accomplish this using a *priority queue*
- Need to be able to read/write bitstrings!
  - Unless we pick multiples of 8 bits for our codewords, we will need to read/write fractions of bytes for our codewords
    - We're not actually going to do I/O on fraction of bytes
    - We'll maintain a buffer of bytes and perform bit processing on this buffer
    - See BinaryStdIn.java and BinaryStdOut.java



# Binary I/O

```
private static void writeBit(boolean bit) {  
    // add bit to buffer  
    buffer <<= 1;  
    if (bit) buffer |= 1;  
    // if buffer is full (8 bits), write out as a single byte  
    N++;  
    if (N == 8) clearBuffer();  
}
```

```
writeBit(true);  
writeBit(false);  
writeBit(true);  
writeBit(false);  
writeBit(false);  
writeBit(false);  
writeBit(false);  
writeBit(true);
```


buffer:

**10100001**

N:

**0**

# Representing tries as bitstrings



# Huffman pseudocode

- Encoding approach:
  - Read input
  - Compute frequencies
  - Build trie/codeword table
  - Write out trie as a bitstring to compressed file
  - Write out character count of input
  - Use table to write out the codeword for each input character
- Decoding approach:
  - Read trie
  - Read character count
  - Use trie to decode bitstring of compressed file

# Further implementation concerns

- To encode/decode, we'll need to read in characters and output codes/read in codes and output characters
  - ...
  - Sounds like we'll need a symbol table!
    - What implementation would be best?
      - Same for encoding and decoding?
  - Note that this means we need access to the trie to expand a compressed file!

# How do we determine character frequencies?

- Option 1: Preprocess the file to be compressed
  - Upside: Ensure that Huffman's algorithm will produce the best output for the given file
  - Downsides:
    - Requires two passes over the input, one to analyze frequencies/build the trie/build the code lookup table, and another to compress the file
    - Trie must be stored with the compressed file, reducing the quality of the compression
      - This especially hurts small files
      - Generally, large files are more amenable to Huffman compression
        - Just because a file is large, however, does not mean that it will compress well!

# How do we determine character frequencies?

- Option 2: Use a static trie
  - Analyze multiple sample files, build a single tree that will be used for all compressions/expansions
  - Saves on trie storage overhead...
  - But in general not a very good approach
    - Different character frequency characteristics of different files means that a code set/trie that works well for one file could work very poorly for another
      - Could even cause an increase in file size after “compression”!

# How do we determine character frequencies?

- Option 3: Adaptive Huffman coding
  - Single pass over the data to construct the codes and compress a file with no background knowledge of the source distribution
  - Not going to really focus on adaptive Huffman in the class, just pointing out that it exists...

# Ok, so how good is Huffman compression

- ASCII requires  $8m$  bits to store  $m$  characters
- For a file containing  $c$  different characters
  - Given Huffman codes  $\{h_0, h_1, h_2, \dots, h_{(c-1)}\}$
  - And frequencies  $\{f_0, f_1, f_2, \dots, f_{(c-1)}\}$
  - Sum from 0 to  $c-1$ :  $\text{len}(h_i) * f_i$
- Total storage depends on the differences in frequencies
  - The bigger the differences, the better the potential for compression
- Huffman is optimal for character-by-character prefix-free encodings
  - Proof in Propositions T and U of Section 5.5 of the text





# That seems like a bit of a caveat...

- Where does Huffman fall short?
  - What about repeated patterns of multiple characters?
    - Consider a file containing:
      - 1000 A's
      - 1000 B's
      - ...
      - 1000 of every ASCII character
    - Will this compress at all with Huffman encoding?
      - Nope!
    - But it seems like it should be compressible...

# Run length encoding

- Could represent the previously mentioned string as:
  - 1000A1000B1000C, etc.
    - Assuming we use 10 bits to represent the number of repeats, and 8 bits to represent the character...
      - 4608 bits needed to store run length encoded file
      - vs. 2048000 bits for input file
      - Huge savings!
- Note that this incredible compression performance is based on a very specific scenario...
  - Run length encoding is not generally effective for most files, as they often lack long runs of repeated characters

# What else can we do to compress files?



# Patterns are compressible, need a general approach

- Huffman used variable-length codewords to represent fixed-length portions of the input...
  - Let's try another approach that uses fixed-length codewords to represent variable-length portions of the input
- Idea: the more characters can be represented in a single codeword, the better the compression
  - Consider "the": 24 bits in ASCII
  - Representing "the" with a single 12 bit codeword cuts the used space in half
    - Similarly, representing longer strings with a 12 bit codeword would mean even better savings!

# How do we know that “the” will be in our file?

- Need to avoid the same problems as the use of a static trie for Huffman encoding...
- So use an adaptive algorithm and build up our patterns and codewords as we go through the file

# LZW compression

- Initialize codebook to all single characters
  - e.g., character maps to its ASCII value
- While !EOF:
  - Match longest prefix in codebook
  - Output codeword
  - Take this longest prefix, add the next character in the file, and add the result to the dictionary with a new codeword

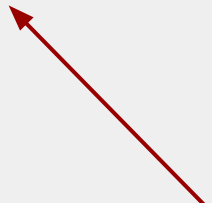
# LZW compression example

- Compress, using 12 bit codewords:
  - TOBEORNOTTOBEORTOBEORNOT

Cur	Output	Add
T	84	TO:257
O	79	OB:258
B	66	BE:259
E	69	EO:260
O	79	OR:261
R	82	RN:262
N	78	NO:263
O	79	OT:264

T	84	TT:265
TO	257	TOB:266
BE	259	BEO:267
OR	261	ORT:268
TOB	266	TOBE:269
EO	260	EOR:270
RN	262	RNO:271
OT	264	--
	256	--

# LZW expansion

- Initialize codebook to all single characters
    - e.g., ASCII value maps to its character
  - While !EOF:
    - Read next codeword from file
    - Lookup corresponding pattern in the codebook
    - Output that pattern
    - Add the previous pattern + the first character of the current pattern to the codebook
- Note this means no codebook addition after first pattern output!
- 



# LZW expansion example

Cur	Output	Add
84	T	--
79	O	257:TO
66	B	258:OB
69	E	259:BE
79	O	260:EO
82	R	261:OR
78	N	262:RN
79	O	263:NO

84	T	264:OT
257	TO	265:TT
259	BE	266:TOB
261	OR	267:BEO
266	TOB	268:ORT
260	EO	269:TOBE
262	RN	270:EOR
264	OT	271:RNO
256		--

# How does this work out?

- Both compression and expansion construct the same codebook!
  - Compression stores character string → codeword
  - Expansion stores codeword → character string
  - They contain the same pairs in the same order
    - Hence, the codebook doesn't need to be stored with the compressed file, saving space

# Just one tiny little issue to sort out...

- Expansion's codebook will always be a step "behind" compression's when processing the same pattern
  - If, during compression, the (pattern, codeword) that was just added to the dictionary is immediately used in the next step, the decompression algorithm will not yet know the codeword.
  - This can be easily detected and dealt with, however

# LZW corner case example

- Compress, using 12 bit codewords: AAAAAA

Cur	Output	Add
A	65	AA:257
AA	257	AAA:258
AAA	258	--

- Expansion:

Cur	Output	Add
65	A	--
257	AA	257:AA
258	AAA	258:AAA

# LZW implementation concerns: codebook

- How to represent/store during:
  - Compression
  - Expansion
- Considerations:
  - What operations are needed?
  - How many of these operations are going to be performed?
- Discuss

# Further implementation issues: codeword size

- How long should codewords be?
  - Use fewer bits:
    - Gives better compression earlier on
    - But, leaves fewer codewords available, which will hamper compression later on
  - Use more bits:
    - Delays actual compression until longer patterns are found due to large codeword size
    - More codewords available means that greater compression gains can be made later on in the process

# Variable width codewords

- This sounds eerily like variable length codewords...
  - Exactly what we set out to avoid!
- Here, we're talking about a different technique
- Example:
  - Start out using 9 bit codewords
  - When codeword 512 is inserted into the codebook, switch to outputting/grabbing 10 bit codewords
  - When codeword 1024 is inserted into the codebook, switch to outputting/grabbing 11 bit codewords...
  - Etc.

# Even further implementation issues: codebook size

- What happens when we run out of codewords?
  - Only  $2^n$  possible codewords for  $n$  bit codes
  - Even using variable width codewords, they can't grow arbitrarily large...
- Two primary options:
  - Stop adding new keywords, use the codebook as it stands
    - Maintains long already established patterns
    - But if the file changes, it will not be compressed as effectively
  - Throw out the codebook and start over from single characters
    - Allows new patterns to be compressed
    - Until new patterns are built up, though, compression will be minimal



# The showdown you've all been waiting for...

## HUFFMAN vs LZW

- In general, LZW will give better compression
  - Also better for compression archived directories of files
    - Why?
      - Very long patterns can be built up, leading to better compression
      - Different files don't "hurt" each other as they did in Huffman
        - Remember our thoughts on using static tries?

# So lossless compression apps use LZW?

- Most dedicated compression applications use other algorithms:
  - DEFLATE (combination of LZ77 and Huffman)
    - Used by PKZIP and gzip
  - Burrows-Wheeler transforms
    - Used by bzip2
  - LZMA
    - Used by 7-zip
  - brotli
    - Published by Google in 2015
  - Zstandard (zstd)
    - Published by Facebook in 2016
    - Designed to provide DEFLATE-like compression ratios with faster expansion runtimes
    - At its maximum compression level gives a compression ratio close to lzma, better than bzip2
    - Currently decompresses faster than any other algorithm while having similar or better compression ratio

# DEFLATE et al achieve even better general compression?

- How much can they compress a file?
- Better question:
  - How much can a file be compressed by any algorithm?
- No algorithm can compress every bitstream
  - Assume we have such an algorithm
  - We can use to compress its own output!
  - And we could keep compressing its output until our compressed file is 0 bits!
    - Clearly this can't work
- Proofs in Proposition 5 of Section 5.5 of the text

# Can we reason about how much a file can be compressed?

- Yes! Using Shannon Entropy



# Information theory in a single slide...

- Founded by Claude Shannon in his paper "A Mathematical Theory of Communication"
- *Entropy* is a key measure in information theory
  - Slightly different from thermodynamic entropy
  - A measure of the unpredictability of information content
  - By losslessly compressing data, we represent the same information in less space
  - Hence, 8 bits of uncompressed text has less entropy than 8 bits of compressed data

# Entropy applied to language:

- Translating a language into binary, the entropy is the average number of bits required to store a letter of the language
- Entropy of a message \* length of message = amount of information contained in that message
- On average, a lossless compression scheme cannot compress a message to have more than 1 bit of information per bit of compressed message
- Uncompressed, English has between 0.6 and 1.3 bits of entropy per character of the message